

HPCSE II - Exercise 1

Anian Ruoss

February 23, 2019

All experiments were run on a node obtained by running the following command:

```
bsub -n 24 -R fullnode -R "select[model=XeonE5_2680v3]" -Is bash
```

Task 4

- a) See the first row in Table 1.
- b) The optimal number of grids is 6 as illustrated in Table 1.

Table 1: Iterations and running times for different numbers of grids for DownRelaxations = 1 and UpRelaxations = 1.

GridCount	Iterations	Running Time	Verification
1	199204	150.165s	Passed
2	43522	57.425s	Passed
3	32457	46.268s	Passed
4	32102	46.524s	Passed
5	30625	44.926s	Passed
6	22973	33.810s	Passed
7	33171	48.768s	Passed
8	51762	75.945s	Passed

- c) The optimal combination of up and down relaxations is determined in a greedy way. The optimal number of down relaxations in terms of running time is 3 as illustrated in Table 2. Although, values of 5 and 7 produce less iterations, the time per iteration is significantly higher.

Table 2: Iterations and running times for different numbers of down relaxations given GridCount = 6 and UpRelaxations = 1.

DownRelaxations	Iterations	Running Time	Verification
1	22973	33.770s	Passed
2	37152	72.471s	Passed
3	13467	33.071s	Passed
4	20885	61.631s	Passed
5	12068	41.526s	Passed
6	14811	58.096s	Passed
7	10130	44.835s	Passed

The optimal number of up relaxations is 1 as illustrated in Table 3.

Table 3: Iterations and running times for different numbers of up relaxations with GridCount = 6 and DownRelaxations = 3.

UpRelaxations	Iterations	Running Time	Verification
1	13467	33.055s	Passed
2	13467	34.066s	Passed
3	13467	35.200s	Passed
4	13467	36.909s	Passed
5	13467	37.349s	Passed
6	13467	38.467s	Passed
7	13467	39.613s	Passed

- d) The optimal configuration is given by GridCount = 6, DownRelaxations = 3, and UpRelaxations = 1.
- e) The solution reaches its equilibrium state faster on coarser grids as every cell covers more area. Fine grids ensure that the discretization error remains small.
- f) TODO

Task 5

- a) *Loop Interchange* can be applied to prevent bad memory access patterns, i.e. swapping the i - and j -loops. *Loop Fusion* can be applied to prevent the overhead of a for-loop, i.e. when two independent computations with the same iteration space are merged into a single for-loop.
- b) Every grid should be allocated in its entirety before moving on to the next one as this prevents cache thrashing.
- c) Omitted.

Task 6

- a)
 - Division operator replaced with multiplication of inverse.
 - Multiplication by 1 removed.
 - $\text{pow}(x, 2)$ replaced with $x \cdot x$.
 - Constants pre-computed and moved out of loops.
 - Expressions simplified with associativity.
 - Copy of array avoided by pointer swap.

Task 7

- a) Most loops cannot be vectorized due to vector dependence.
- b) This can be fixed by adding `#pragma ivdep` before the loop if the vectors are in fact independent.
- c) Unaligned SIMD operations need to gather from different locations, whereas aligned operations rely on the fact that the relevant data is stored contiguously in memory.

Table 4: Iterations and running times for the three problems after all optimizations have been performed.

Problem	Iterations	Running Time	Verification
1	7844	1.422s	Passed
2	13467	8.238s	Passed
3	12865	7.824s	Passed